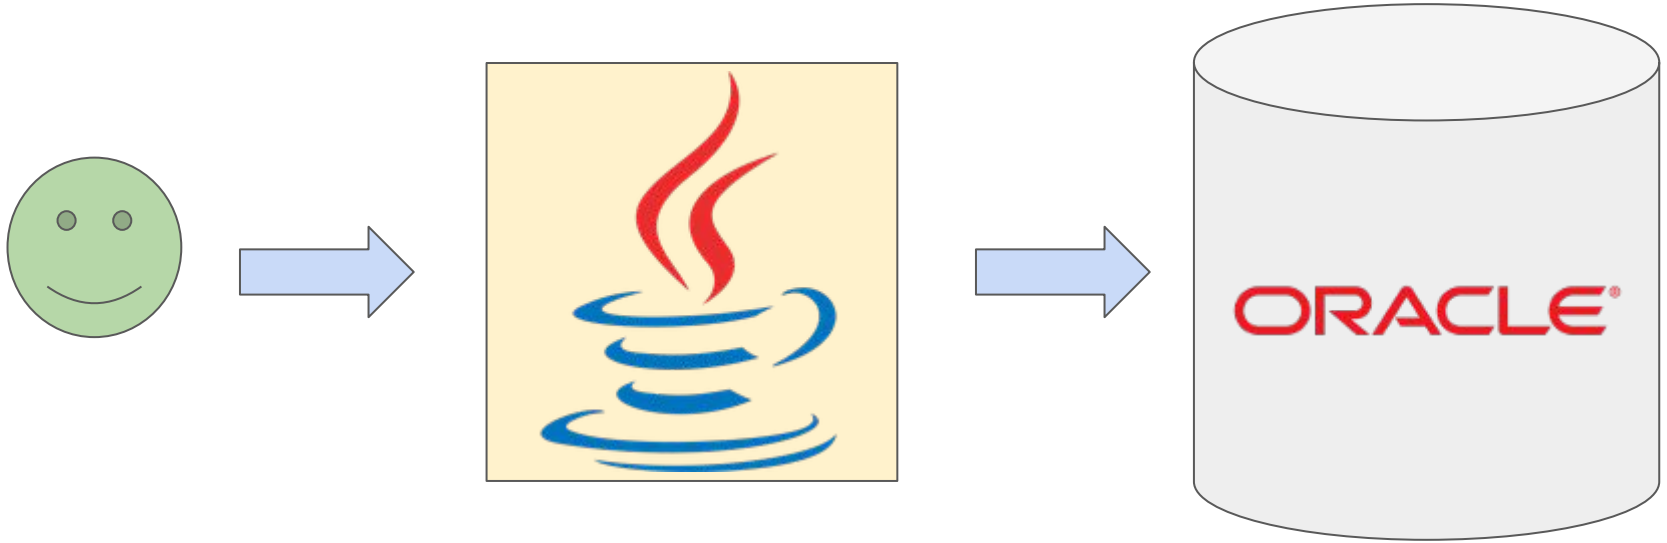




monitors distributed systems

*a long time ago in a galaxy far far away...*



# Cloud Native Landscape

v1.0

App Definition & Development

Database & Data Analytics

Streaming

SCM

Application Definition

CI/CD

Orchestration & Management

Scheduling & Orchestration

Coordination & Service Discovery

Service Management

Runtime

Cloud-Native Storage

Container Runtime

Cloud-Native Network

Provisioning

Host Management / Tooling

Infrastructure Automation

Container Registries

Secure Images

Key Management

Cloud

Public

Private

Platforms

PaaS / Container Service


Serverless/Event-based

Observability & Analysis


Monitoring



Logging

Tracing

  
github.com/cncf/landscape

This landscape is intended as a map through the previously uncharted terrain of cloud native technologies. There are many routes to deploying a cloud native application, with CNCF Projects representing a particularly well-traveled path.

  
**CLOUD NATIVE COMPUTING FOUNDATION**



Grayed out are open source projects

**Distributed architectures are hard**

# Monitoring distributed systems

100k write/sec



*cassandra*

105k write/sec



*cassandra*

101k write/sec

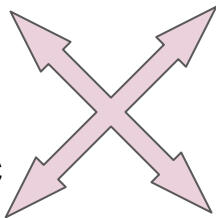


*cassandra*

1k write/sec



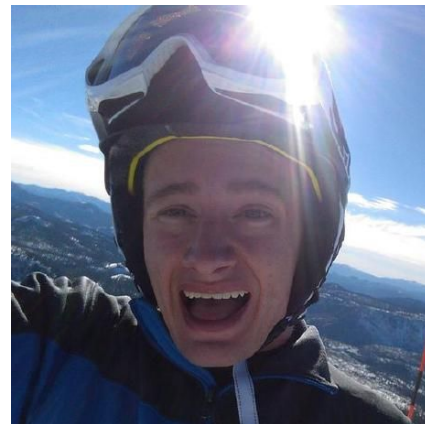
*cassandra*



- Time windows
- Rates
- Percentiles
- Cluster monitoring
- Correlation between metrics
- State transitions (OK => KO)
- Alerts (mail, slack, pagerduty...)
- **Flexibility**
- ...

**RIEMANN**

- Created by Kyle Kingsbury (Aphyr)
- Event processing
- Clojure
- Monitoring



# An immutable event



A diagram showing an immutable event as a collection of key-value pairs. The keys are on the left, and the values are on the right. A large left curly brace groups all the keys, and a large right curly brace groups all the values. The pairs are:

<b>:host</b>	<b>“foo.bar.com”</b>
<b>:service</b>	<b>“df_percent_bytes_used_root”</b>
<b>:state</b>	<b>“critical”</b>
<b>:time</b>	<b>1493243041</b>
<b>:metric</b>	<b>90</b>
<b>:description</b>	<b>“Disk is full”</b>
<b>:tags</b>	<b>[“disk”]</b>
<b>:ttl</b>	<b>60</b>

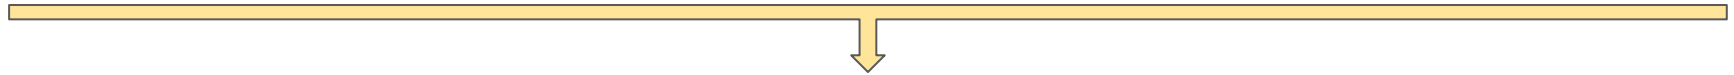


Collectd  
Telegraf  
K8s/Heapster  
Statsd  
Graphite  
...

Syslog-ng  
Logstash  
Fluentd  
...

Java  
Haskell  
Go  
Python  
Perl  
...

Kafka  
Nagios check  
Chef  
...



**Good**

TCP  
TLS

**Drop packets**

UDP

**Slow**

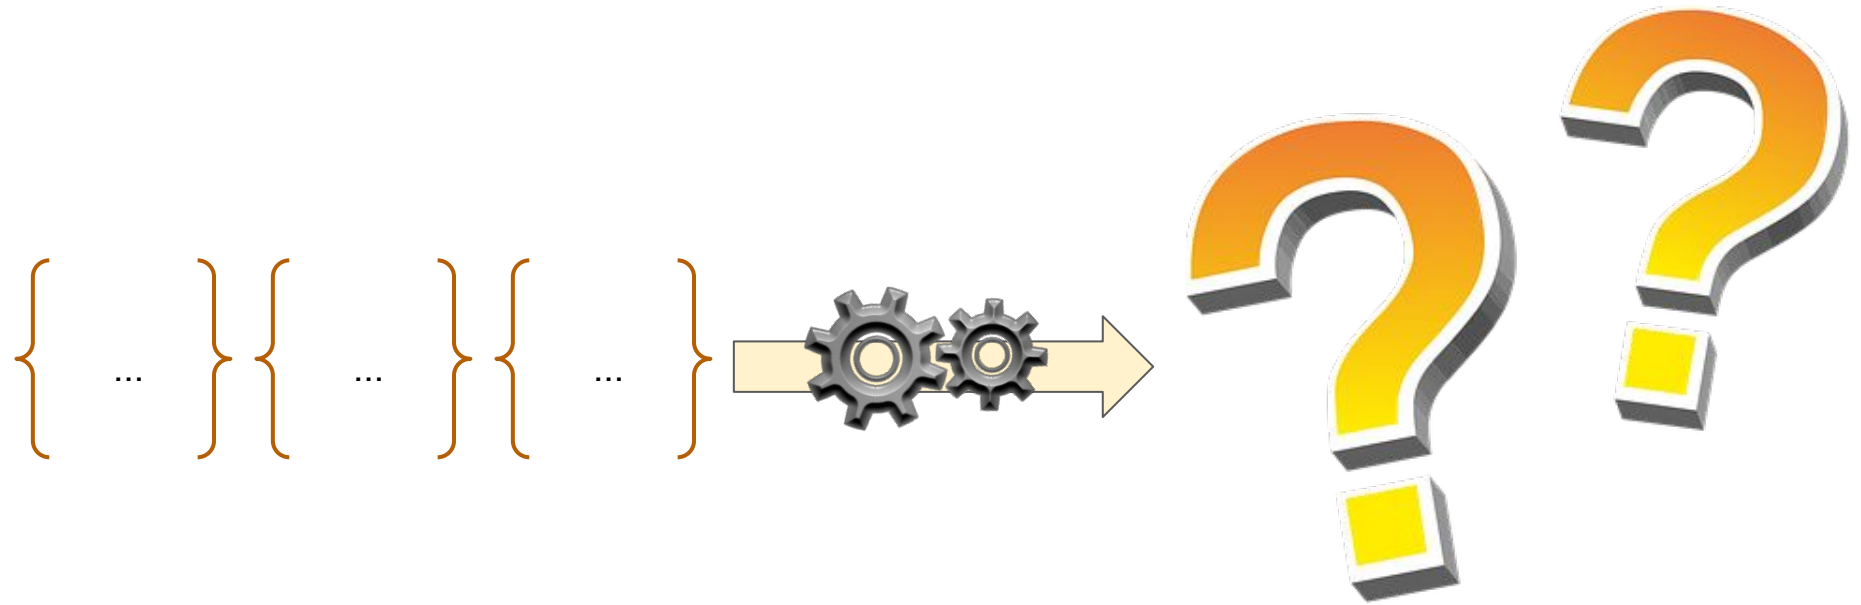
HTTP

**Compat**

Graphite  
OpenTSDB

**RIEMANN**

# Streams



:host	"foo1.com"	:host	"foo1.com"	:host	"foo1.com"
:service	"api_rate"	:service	"foobar"	:service	"api_rate"
:time	1493243041	:time	1493243041	:time	1493243044
:metric	90	:metric	90	:metric	90



**where** = service "api\_rate"

:host	"foo1.com"
:service	"api_rate"
:time	1493243041
:metric	90

:host	"foo1.com"
:service	"api_rate"
:time	1493243044
:metric	90

**:host** "foo1.com"  
**:service** "api\_rate"  
**:time** 1493243041  
**:metric** 90

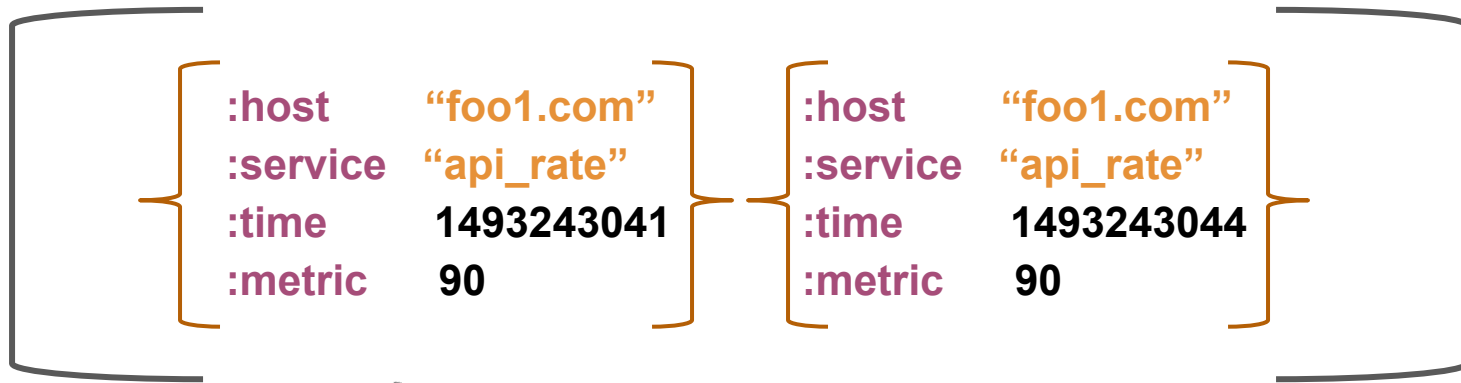
**:host** "foo1.com"  
**:service** "api\_rate"  
**:time** 1493243044  
**:metric** 90



**fixed-time-window 10**

**:host** "foo1.com"  
**:service** "api\_rate"  
**:time** 1493243041  
**:metric** 90

**:host** "foo1.com"  
**:service** "api\_rate"  
**:time** 1493243044  
**:metric** 90



**smap** sum



:host	"foo1.com"
:service	"api_rate"
:time	1493243044
:metric	180



**where** < metric 200

:host	"foo1.com"
:service	"api_rate"
:time	1493243044
:metric	180

:host	"foo1.com"
:service	"api_rate"
:time	1493243044
:metric	180



**email** "ops@riemann.io"





**where** = service "api\_rate"



**fixed-time-window** 10



**smap** sum



**where** < metric 200



**email** "ops@riemann.io"





(**where** (= service “api\_rate”)



(**fixed-time-window** 10



(**smap** sum



(**where** (< metric 200)



(**email** “ops@riemann.io”))))))



(fi



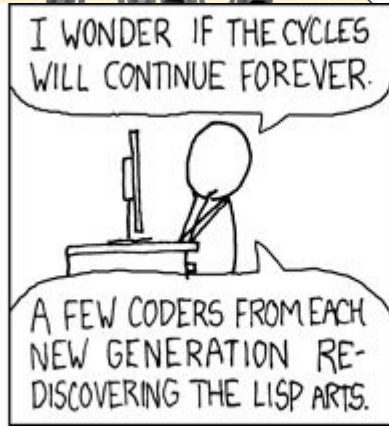
pi\_rate")

10



Use map !

```
riemann.bin> "REPL driven development ftw"  
"REPL driven development ftw"
```



# Configuration as code

(Your config is 100 % Clojure)

```
(with {:description "Disk is full"  
      :state        "critical"}  
child)
```

{  
:host "foo.bar.com"  
:service "df\_home\_mathieu"  
:state "ok"  
:time 1493243041  
:metric 90  
}



{  
:host "foo.bar.com"  
:service "df\_home\_mathieu"  
:state "ok"  
:time 1493243041  
:metric 90  
:description "Disk is full"  
}

```
(where (service "foo")
```

**where** has  
2 children

```
  (with {:description "cat"}  
    (email "ops@riemann.io"))
```

First child

```
  (with {:description "dog"}  
    (email "dev@riemann.io")))
```

Second child

```
(where (service "foo")
```

**where** has  
2 children

**with** has  
1 child

```
(with {:description "cat"}  
  (email "ops@riemann.io"))
```

First child

**with** has  
1 child

```
(with {:description "dog"}  
  (email "dev@riemann.io"))
```

Second child

# Clojure datastructures

## Immutability

### No side effects between streams

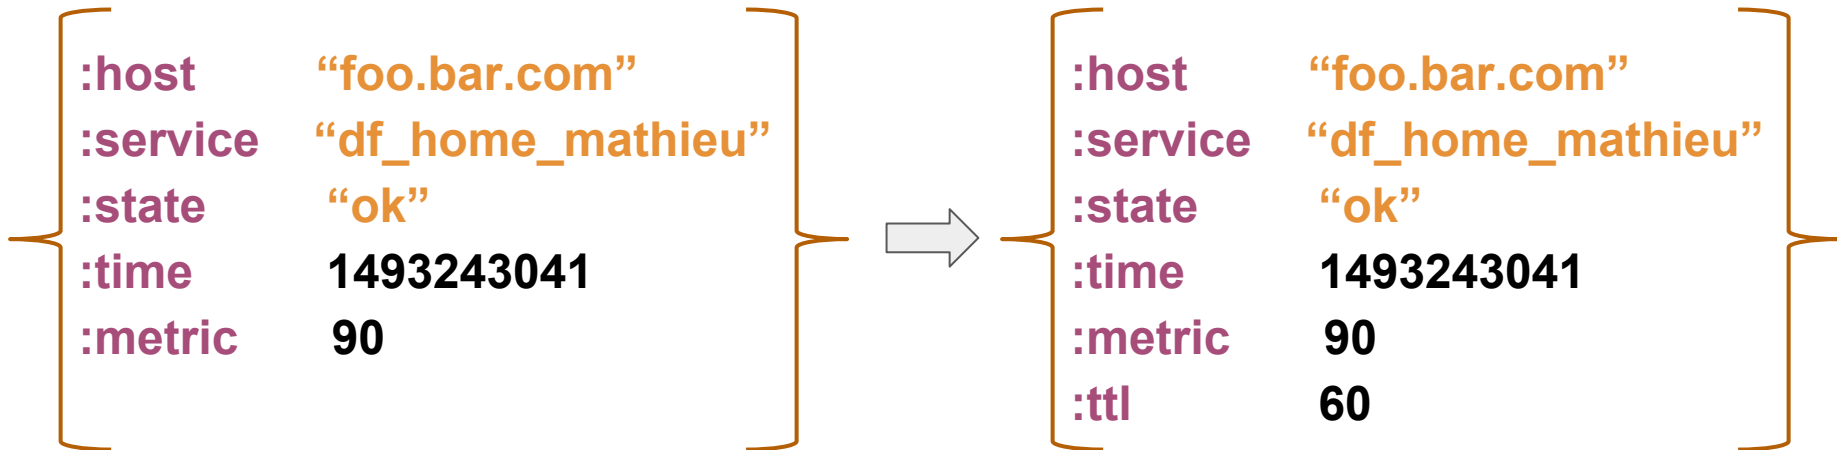
```
(where (service "df_percent_bytes_used_var_log"))
```

```
(where (service #"^df_percent_bytes_used_"))
```

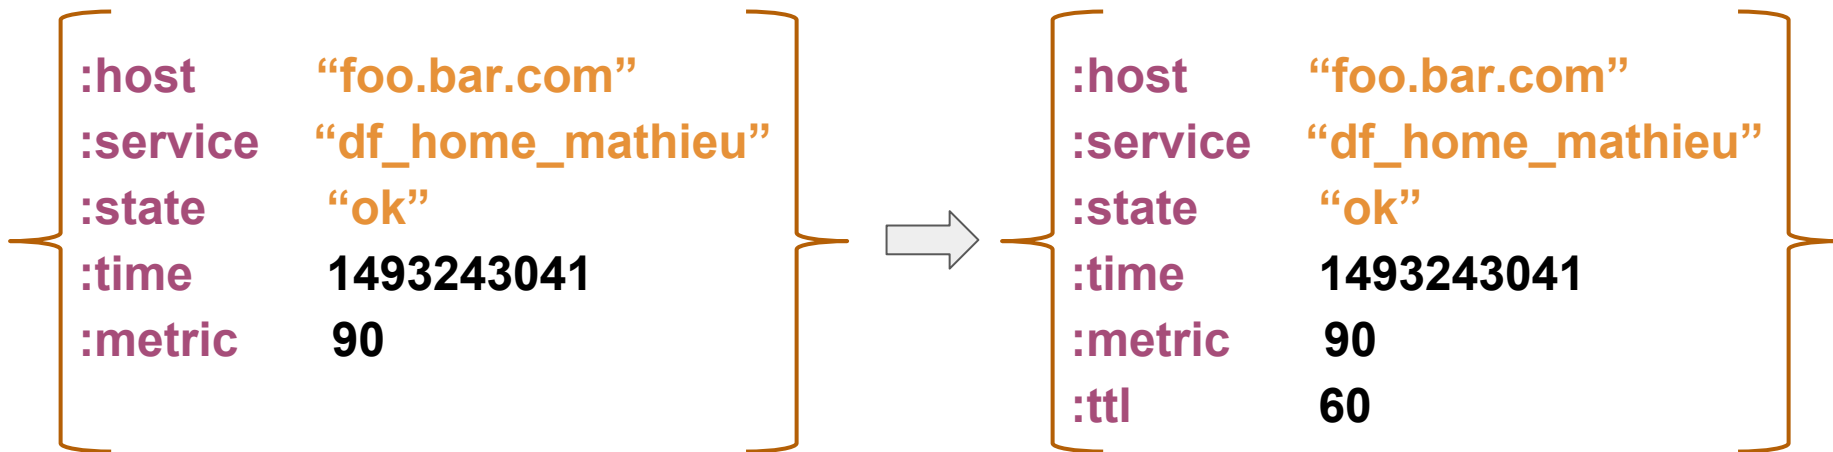
```
(where (and (service #"^df_percent_bytes_used_")  
             (> (:metric event) 80)))
```



(default :ttl 60  
child)



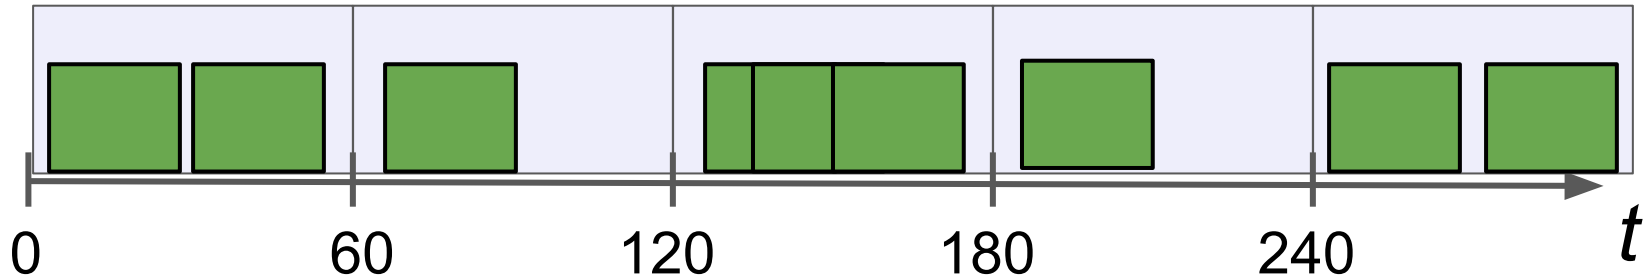
```
(smap (fn [event]  
  (assoc event :ttl 60))  
child)
```



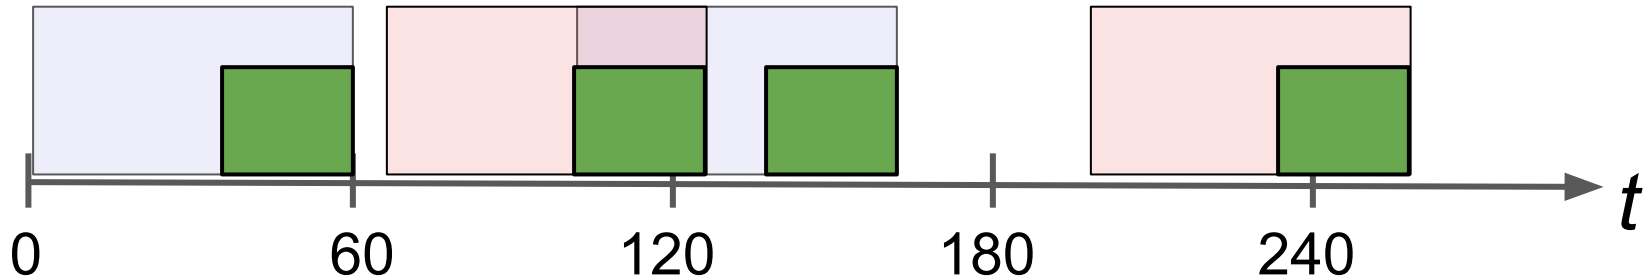
(fixed-time-window 60

child1

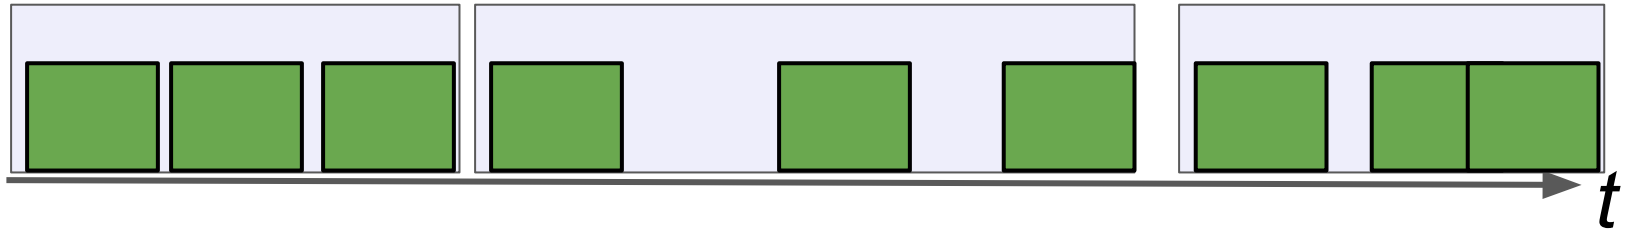
child2)



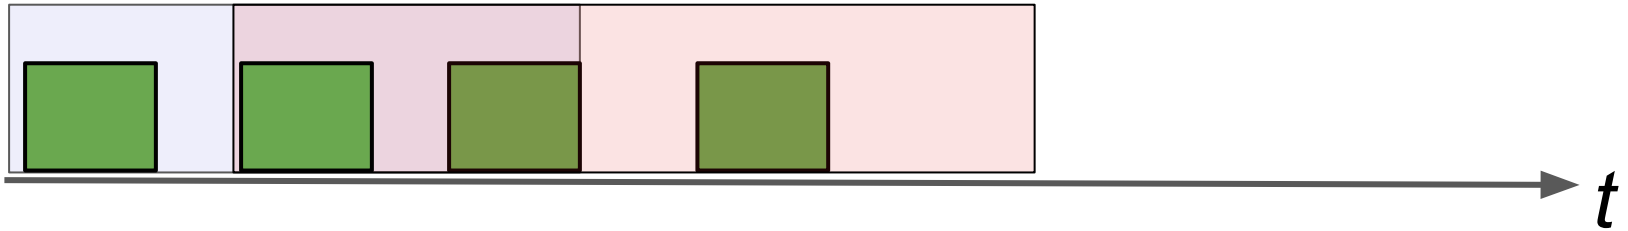
(moving-time-window 60  
child)



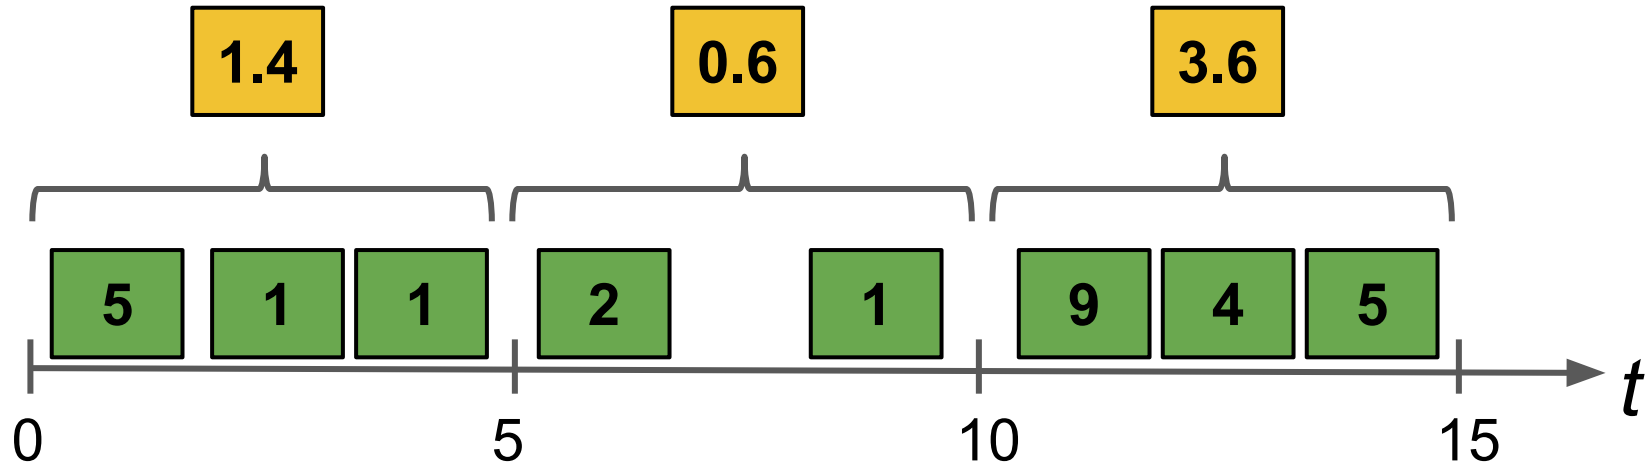
(fixed-event-window 3  
child)



(moving-event-window 3  
child)

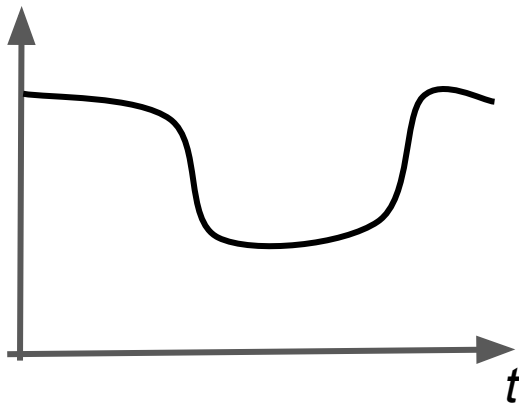


(rate 5  
child)

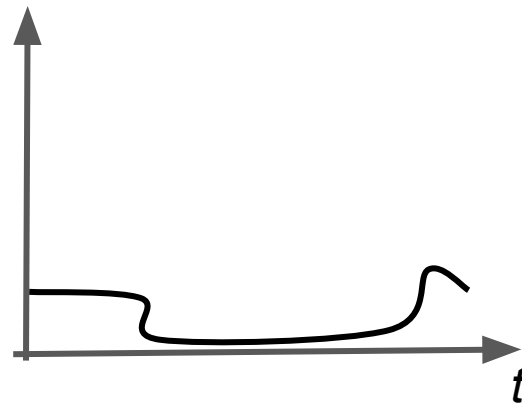


(scale (/ 1 1024 1024 1024)  
child)

*bytes*

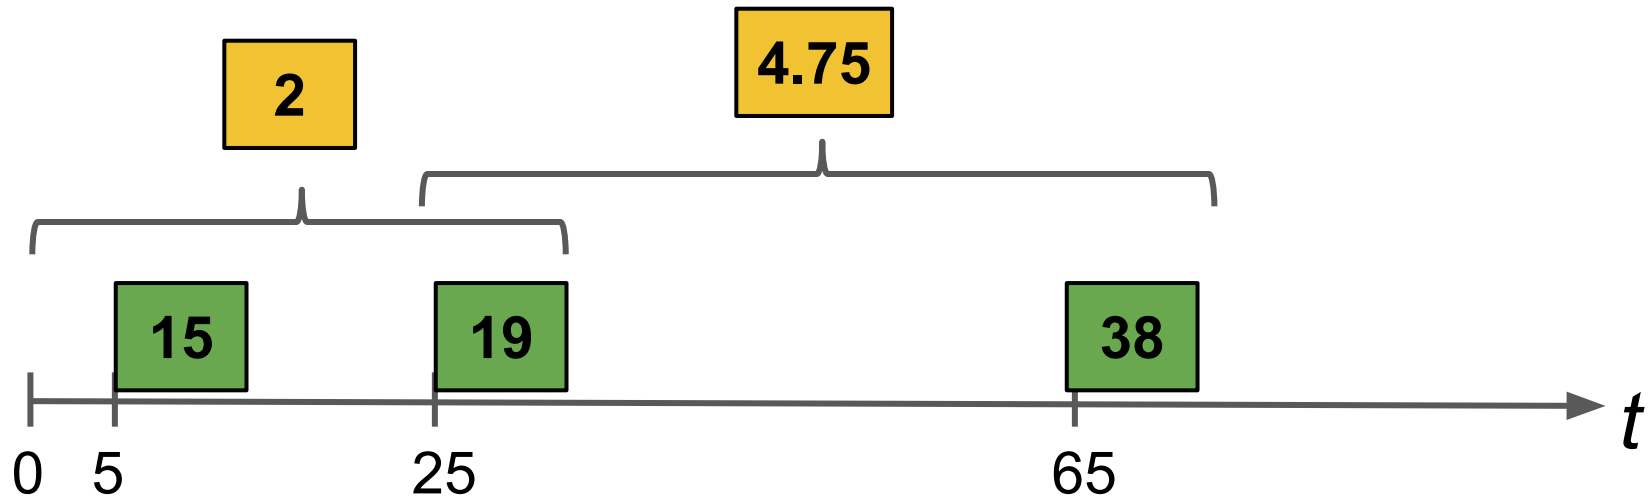


*Gigabytes*





(ddt  
child)



```
(changed :state {:init "ok"}  
child)
```

Sent  
to child



ko

ok

ko

ok

ok

ko

ko

ko

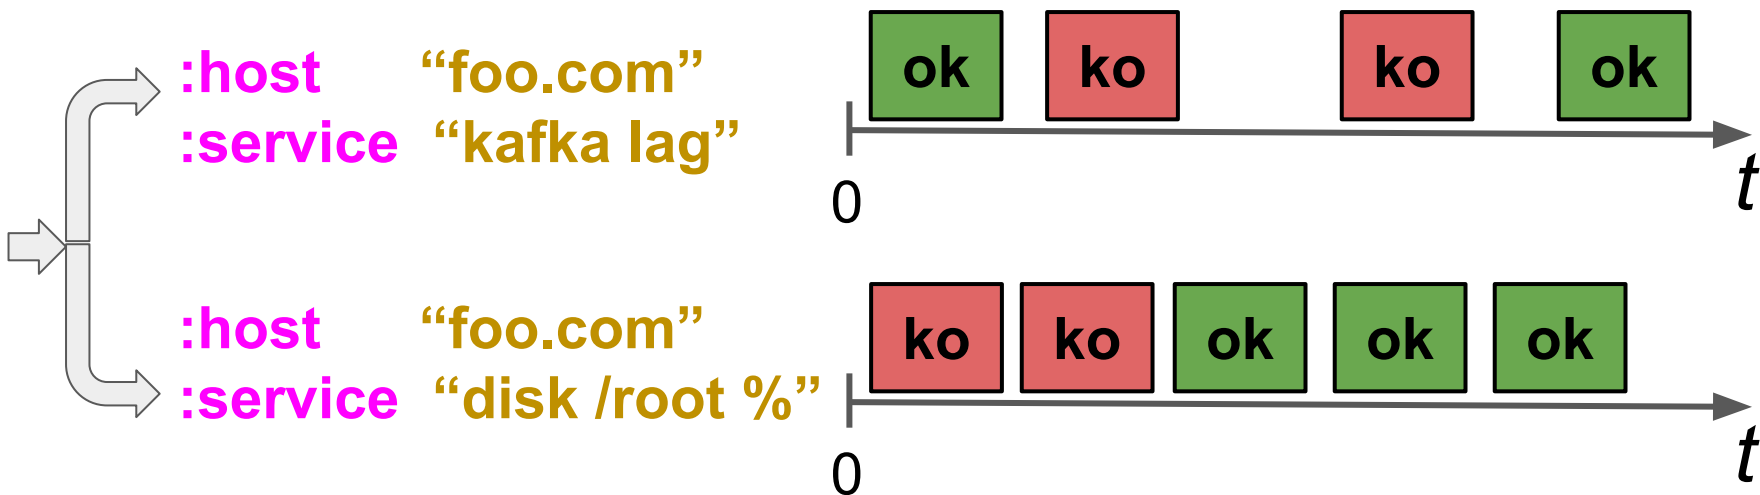
ok

ko

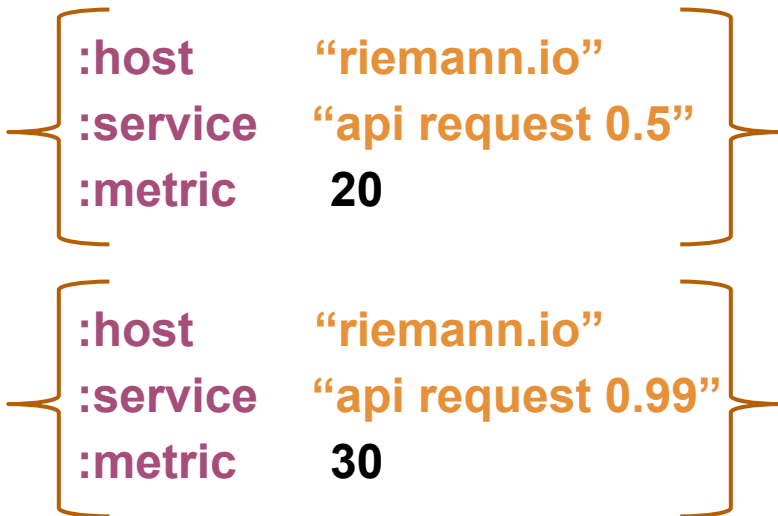
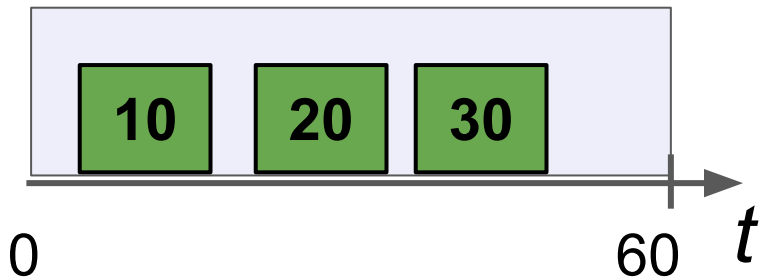
0

$t$

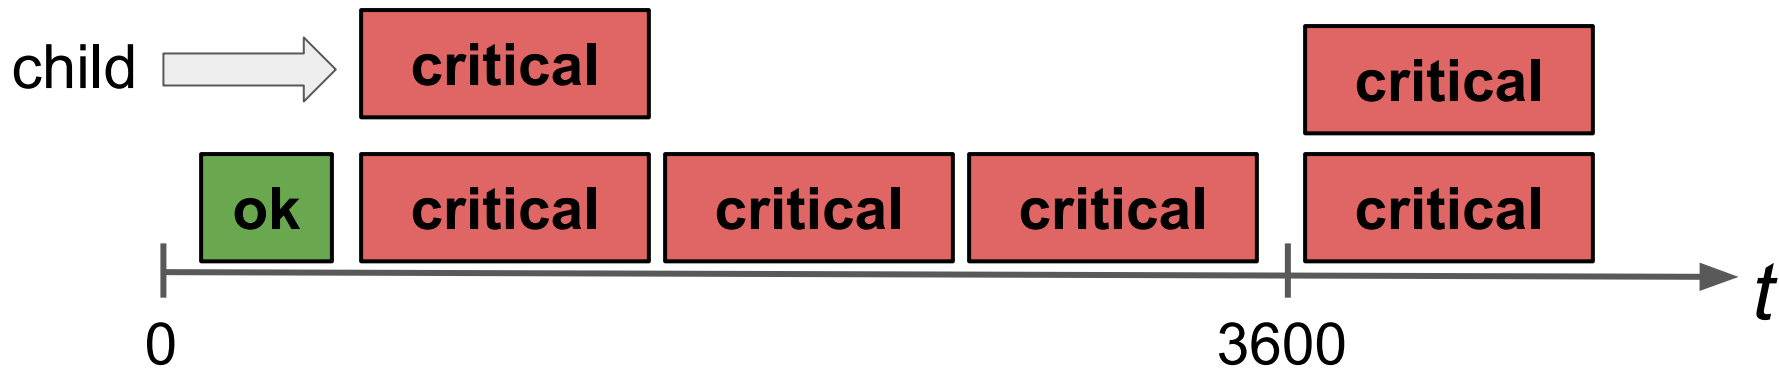
```
(by [:host :service]
  (changed :state {:init "ok"}
    child))
```



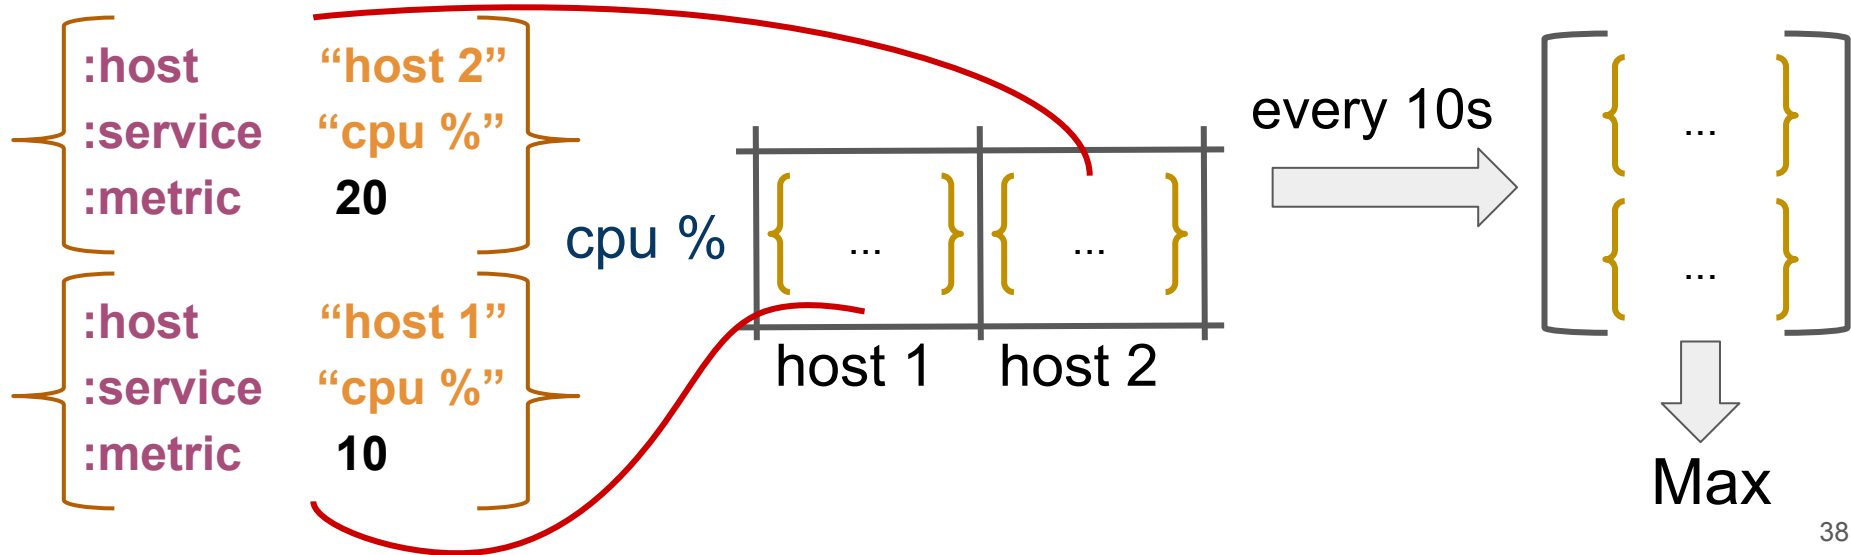
```
(where (service "api request")  
  (percentiles 60 [0.5 0.99]  
    child))
```



```
(where (state "critical")  
  (throttle 2 3600  
    (email "foo@riemann.io"))))
```



(where (service “cpu %”)  
(coalesce 10  
(smap max)))



# RIEMANN

InfluxDB

Elasticsearch

Graphite

Kafka

Logstash

Datadog

Cloudwatch

Riemann

...

Email

Hipchat

Slack

Mailgun

...

Pagerduty

VictorOps

Twilio

Alerta

...

Nagios

Shinken

...

```
(batch 100 1 ;; batch size = 100 every 1 sec
(async-queue! :influxdb ;; create a threadpool
{:queue-size 10000
:core-pool-size 4}
(influxdb {:host 127.0.0.1 ;; forward to influx
:db "riemann"})))
```



```
(exception-stream  
  (email "alert@riemann.io")  
  (influxdb { :host 127.0.0.1  
              :db "riemann" })))
```

Configuration as code

**Split your configuration**

```
(def check-critical-state      ;; a var containing a stream  
  (where (state "critical")  
    (email "admin@riemann.io")))
```

```
(defn check-state      ;; a function returning a stream  
  [s email-addr]  
  (where (state s)  
    (email email-addr)))
```

```
(streams  
  check-critical-state  
  (check-state "critical" "admin@riemann.io"))
```

/etc/riemann/riemann.config

/mycorp/app/elasticsearch.clj

/mycorp/output/mail.clj

/mycorp/system/disk.clj

/mycorp/system/ram.clj

...

+ A plugin system

Configuration as code

**Tests**

```
(scale (/ 1 1024 1024 1024)
  (tap :scale-tap)
  child)
```

```
(tests
  (deftest foo-test
    (is (= (:scale-tap (inject! [{:metric 1000}]))
      [{:metric (/ 1000 1024 1024 1024)}]))))
```

# The index

- In memory datastructure (hashmap)
  - Key : [host service]
  - Value: an event
  - The index stream adds event to the index

```
(where (service "ram_percent")  
  (index))
```

service

Time  
3

ram\_%

{  
:host "foo.bar"  
:service "ram\_percent"  
:metric 65  
:ttl 120  
:time 2  
}

{  
:host "fizz.buzz"  
:service "ram\_percent"  
:metric 80  
:ttl 120  
:time 2  
}

cpu\_%

{  
:host "foo.bar"  
:service "cpu\_percent"  
:metric 40  
:ttl 60  
:time 1  
}

{  
:host "fizz.buzz"  
:service "cpu\_percent"  
:metric 90  
:ttl 60  
:time 3  
}

foo.bar

fizz.buzz

host



service

Time  
10

ram\_%

{  
:host "foo.bar"  
:service "ram\_percent"  
:metric 65  
:ttl 120  
:time 2  
}

{  
:host "fizz.buzz"  
:service "ram\_percent"  
:metric 80  
:ttl 120  
:time 2  
}

cpu\_%

{  
:host "foo.bar"  
:service "cpu\_percent"  
:metric 45  
:ttl 60  
:time 10  
}

{  
:host "fizz.buzz"  
:service "cpu\_percent"  
:metric 90  
:ttl 60  
:time 3  
}

foo.bar

fizz.buzz

host

service

Time  
64

ram\_%

{  
:host "foo.bar"  
:service "ram\_percent"  
:metric 65  
:ttl 120  
:time 2  
}

{  
:host "fizz.buzz"  
:service "ram\_percent"  
:metric 80  
:ttl 120  
:time 2  
}

cpu\_%

{  
:host "foo.bar"  
:service "cpu\_percent"  
:metric 45  
:ttl 60  
:time 10  
}

{  
:host "fizz.buzz"  
:service "cpu\_percent"  
:metric 90  
:ttl 60  
:time 3  
}

foo.bar

fizz.buzz

host

service

Time  
64

ram\_%

**:host** "foo.bar"  
**:service** "ram\_percent"  
**:metric** 65  
**:ttl** 120  
**:time** 2

**:host** "fizz.buzz"  
**:service** "ram\_percent"  
**:metric** 80  
**:ttl** 120  
**:time** 2

cpu\_%

**:host** "foo.bar"  
**:service** "cpu\_percent"  
**:metric** 45  
**:ttl** 60  
**:time** 10

**:host** "fizz.buzz"  
**:service** "cpu\_percent"  
**:metric** 90  
**:ttl** 60  
**:time** 3

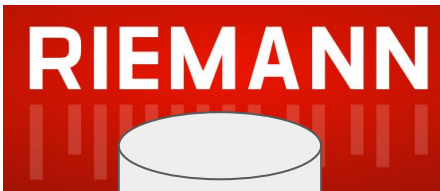
Reinjected in  
Riemann with  
**:state** "expired"

foo.bar

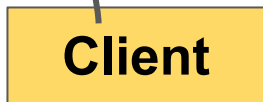
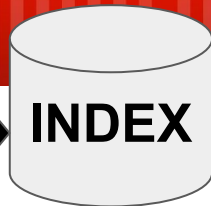
fizz.buzz

host

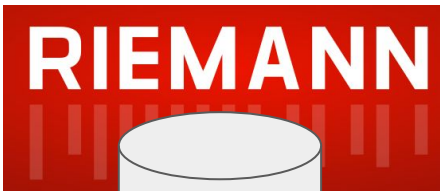
```
(expired  
  (email "expired@riemann.io"))
```



Query =  
"state = cpu\_percent"

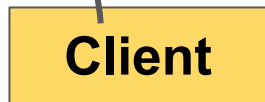
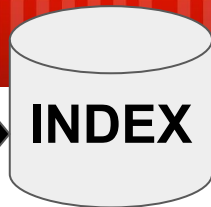


Response =  $\left( \left\{ \dots \right\} \left\{ \dots \right\} \left\{ \dots \right\} \right)$

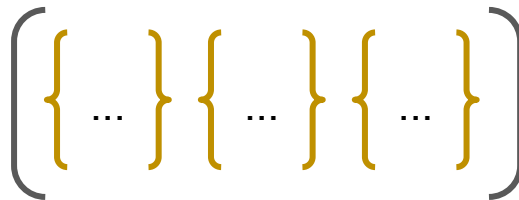


Query =  
"state = cpu\_percent"

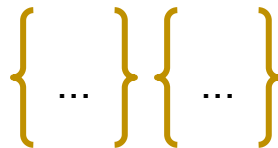
Websocket / SSE  
Stream of events

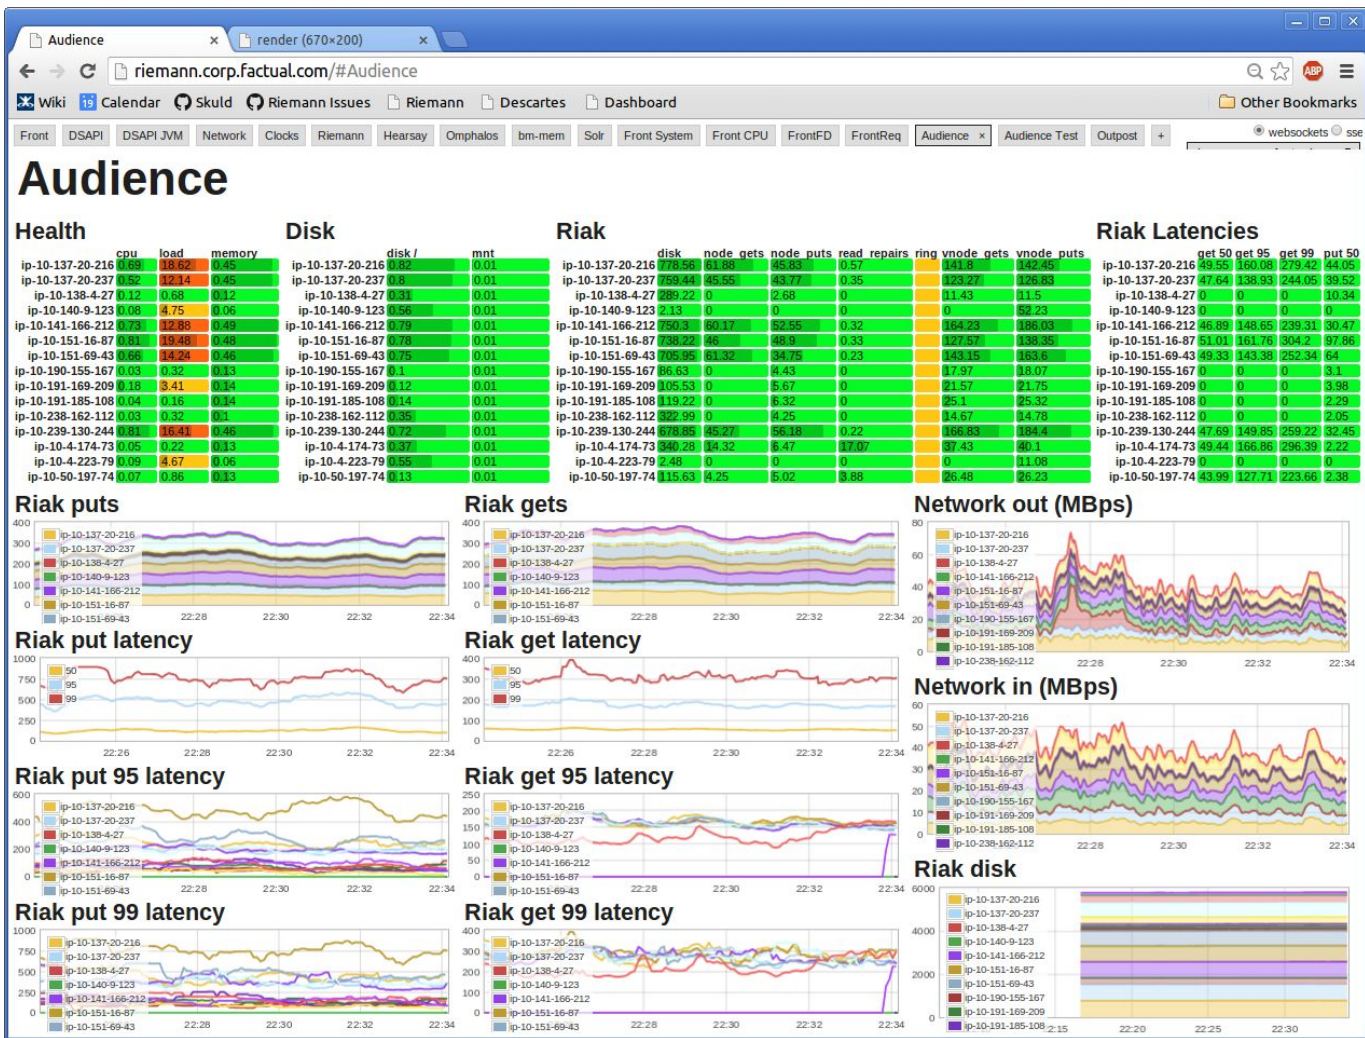


Response =



Query =  
"state = ram\_percent  
and host = fizz.buzz"





JVM



Multithreaded (<3 Clojure)

Netty

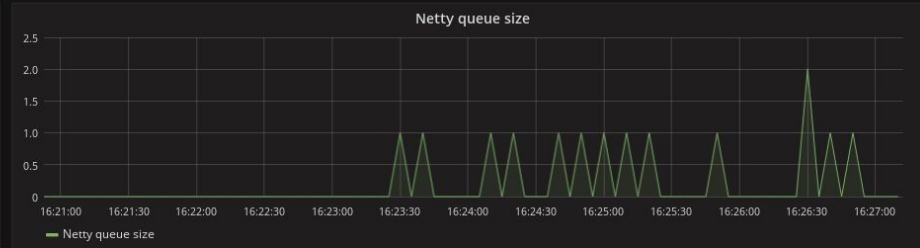
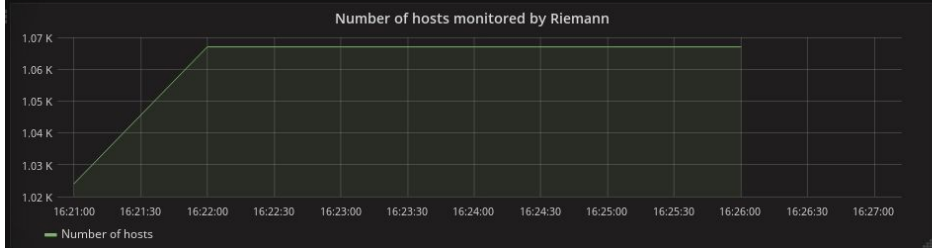
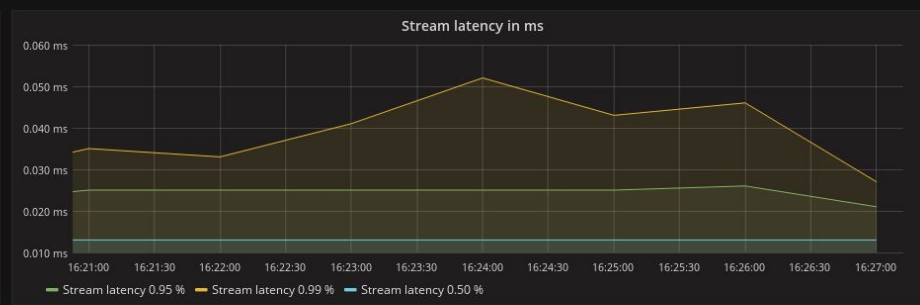
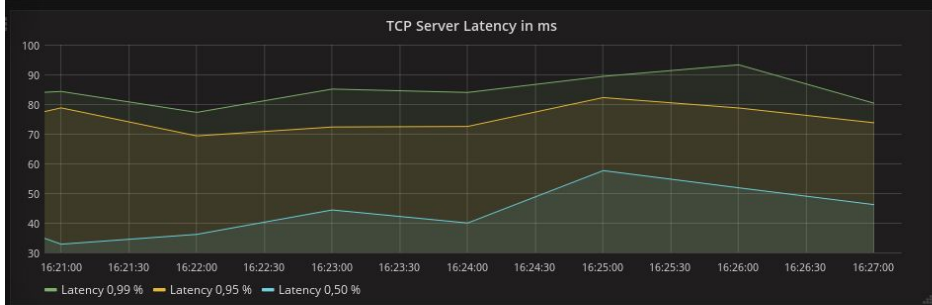
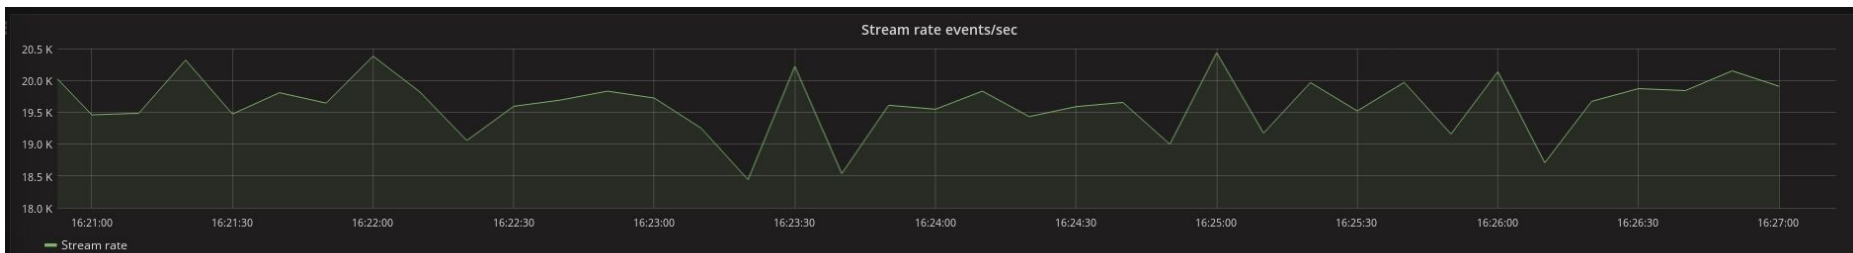
Protobuf

In Memory

Back pressure

...

**FAST**



# No HA

- Sharding
- Send events to 2 instances
- Keepalived

But Clojure is hard ! (((((lisp))))))

## Prometheus :

```
ALERT IncreasedErrorRate
IF (sum(backend_code:haproxy_server_http_responses_total:irate1m{tier="lb", environment="prd", code="5xx"}) by (code,backend)) > 5
FOR 15s
LABELS {severity="critical"}
ANNOTATIONS {
    title="Increased Error Rate Across Fleet",
    description="We are having a high rate of 5xx accross the fleet. It's likely that customers are impacted.",
    runbook="troubleshooting/high-error-rate.md"
}
```

## Zabbix :

```
{Template PfSense:hrStorageFree[{#SNMPVALUE}].last()}<{Template PfSense:hrStorageSize[{#SNMPVALUE}].last()}*0.1
```

But Clojure is hard ! (((((lisp))))))



Thanks !

**riemann.io**

Questions ?